

TASK EXECUTION

Long Horizon Tasks



SOURCE STACK

1K–500K Lines · 6-30 Hours




TARGET STACK



HIDDEN EVALUATION

1 INTEGRITY GATE



Did the migration actually happen?

ANY FAIL → 0


2 FUNCTIONAL TEST



Does it still behave the same?

6,000+ Tests Per Task

3 AGENTIC VERIFIER



Anything still broken?

6 Models × 1 Hour

A near-miss is not a near-pass — one surviving bug puts the whole rewrite in question.