

TASK EXECUTION

Long Horizon Tasks



SOURCE STACK

1K–500K Lines · 6-30 Hours



TARGET STACK



HIDDEN EVALUATION

1 MIGRATION AUDIT



Did the migration
actually happen?

ANY FAIL → 0


2 BEHAVIOURAL TEST



Does it still behave
the same?

6,000+ Tests Per Task

3 AGENTIC VERIFICATION



Anything still broken?

6 Models × 1 Hour

A near-miss is not a near-pass — one surviving bug puts the whole rewrite in question.